

Creating your CUJ (critical User Journey) Lo-Fi

What does your application does and how does it deliver customer value

These are your customers / users

Defining the outlines of your CUJ

Defining the boundaries and patterns of your CUJ

Defining the SLIs

Defining the SLOs

Defining the failure modes

Who are your customers?

what are the entry and exit points of your app?

what are your dependencies

what are you delivering to your users

What are we committing to with those deliverables

How could things go wrong?

other apps

entry

dependency on you to deliver

eg. inability to validate if an account has enough funds to complete a payment request

Other teams

exit

dependency on someone else to deliver to you

end users

Who we are and how can we work with you?

We are Payments Observability

In the pursuit of a metrics-driven approach, we're partnering with payments teams to **accelerate your DevSecOps maturity** by establishing a **metrics-driven observability model**.

Our goal is for all payments teams to have **aligned alerts based on clear Service Level Indicators (SLIs) and Service Level Objectives (SLOs)**.

This ensures our monitoring is targeted, proactive, and directly supports the reliable delivery of critical payment services.

clearer on what is their starting point

show the example dashboard

show a step by step plan they can follow

pager duty is configurable

its made up by the support groups

there is no coding, its all configuration

rolling the data in to SLIs and SLOs is also all configuration

its quick and Iterable

how do we start customer matrix in to measurable metrics?

is this a separate piece of work separate to the technical aspect

is this in addition to our technical scope?

is what we are doing in this CUJ more about the customer scope

we are starting with the customer side but we can definitely support on the technical journey if that is where you are on

we are trying to move away from logs and moving towards basing observability on metrics

and have the ability to consolidate the metric in to something that can show something tangible about what the application is performing

products have multiple application metrics that are going on

how do we measure the metrics with respect of the products when the products have lots of applications involved

we are trying to pivot away from what the memory usage is

we are trying to move towards how do we move in to delivering value

if an application is generic metrics and the metrics are going in to obstack we can draw the broader flows

Critical User Journey (CUJ) Guide

This guide helps you define at a high level how your application delivers value to its users by mapping out key interactions and dependencies.

Define Your Customers/Users

Start by identifying who your application serves, this could be other applications, Internal teams or End users.

Outline your failure modes

Identify potential risks: What could go wrong in this journey? What does failure look like in terms of delivering to your users?

Define the outlines of a Journey

Based on a key user/s that you outlined earlier - Map the flow/journey from start to finish:

Entry point: Where does the user or system begin interacting with your app?

Exit point: What marks the completion of the journey?

Define SLIs (Service Level Indicators)

Determine what you'll measure to track performance: Choose indicators that reflect the health of the CUJ (e.g., latency, error rate, availability).

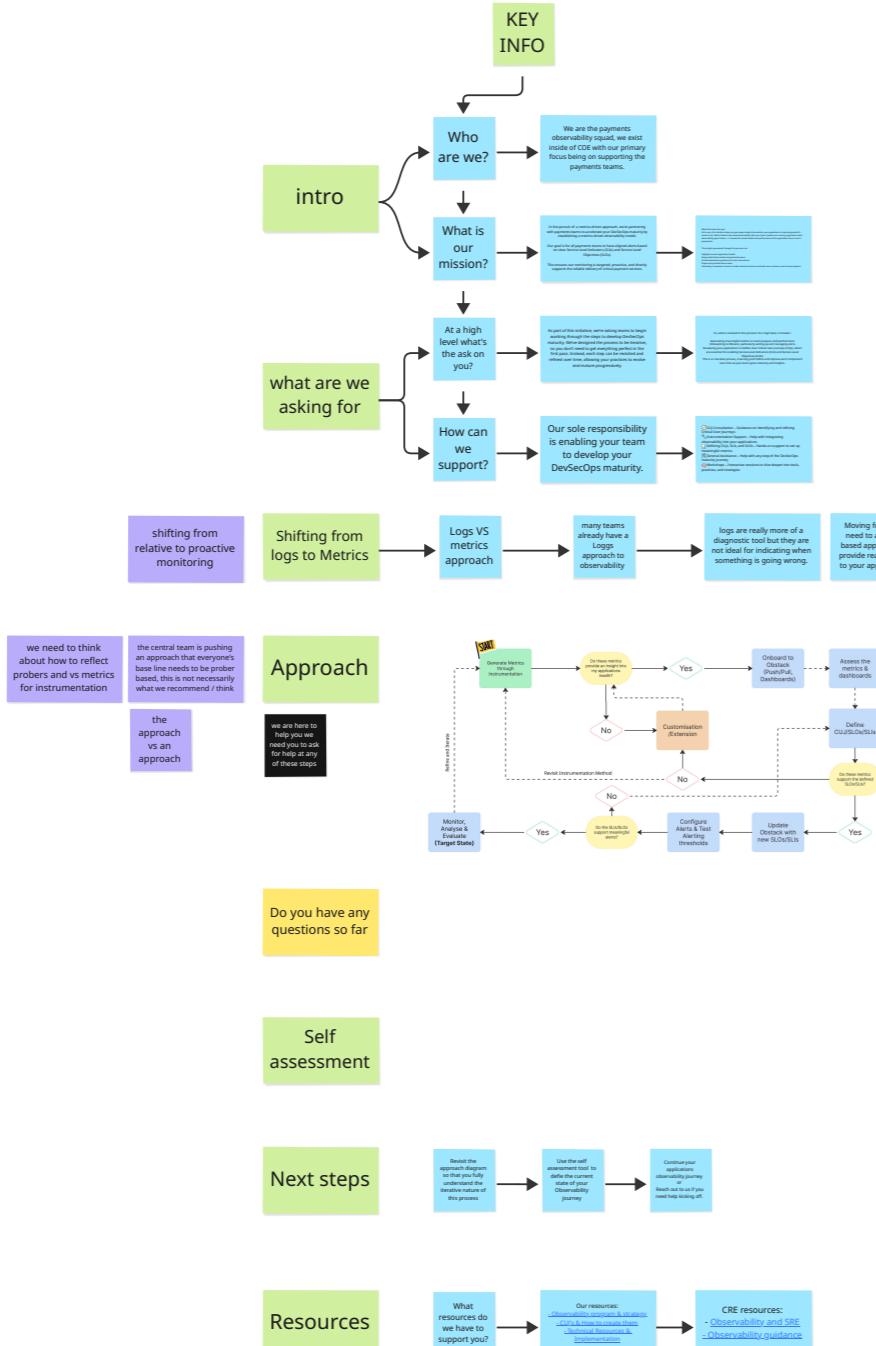
Boundaries & Dependencies

Once you've outlined a complete end-to-end flow that your application participates in, identify and define the segments where your application operates independently - those parts that are not reliant on external systems or integrations, and are solely handled within your application's boundaries

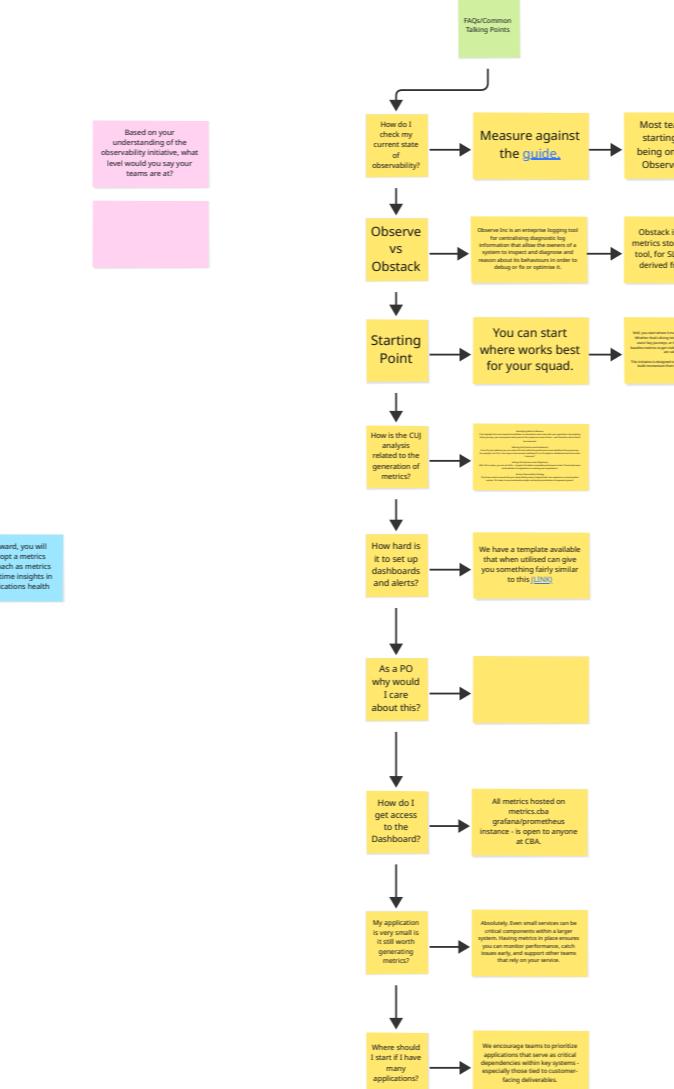
Define SLOs (Service Level Objectives)

Set performance targets: What level of service is acceptable? Use these to guide reliability goals and team priorities.

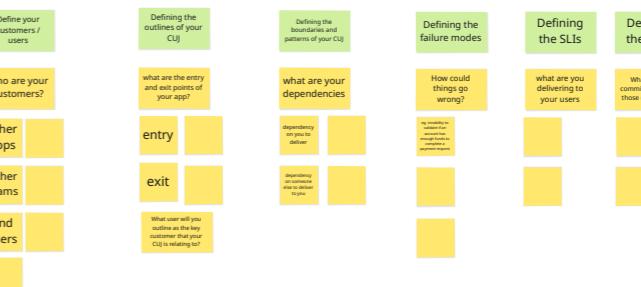
Agenda from intro sessions



Capability measure	Level 0	Awareness Level 1	Enablement Level 2	Embedding Level 3
What type of observability data exists for your technical environment?	The data is support diagnostic data, but it is not readily available.	Existing logging, metrics, and tracing data is available, but may not be fully integrated or require manual effort to access.	Documentation and tools for the different levels of observability are available, and responses to incidents are quick.	All observability data for the different levels of observability are available, and responses to incidents are fast and efficient.
How does your squad find out that your probes are broken?	Failure of the service is communicated via email or alerts are referred to a specific person who is responsible for the engineering work.	Alerts capturing failures are sent to the squad responsible for the probe, but may not be fully integrated with the engineering work.	Alerts are received before the probe fails, and there is a degraded observable service for the probe.	Proactive failure detection and recovery mechanisms are in place, and alerts are received quickly.
How does your squad measure the reliability of your service?	The reliability of the service is measured by looking at logs, but no specific reports or dashboards are published with long term data.	Reliability of the service is measured by reviewing logs and dashboards, but no specific reports or dashboards are published with long term data.	Reliability functions with SLIs and SLUs are integrated with the Service Level Objectives (SLOs) and Service Level Agreements (SLAs). The reliability is measured by feedback in SLO/governance.	Service Level Objective (SLO) and Service Level Agreement (SLA) are well defined and integrated with the reliability of the service.



Creating your CUJ (critical User Journey) Lo-Fi
What does your application do and how does it deliver customer value



Each level is a different pass of the iterative process. Level 3 demands constant evaluation of the alerts, SLIs, SLOs and continuous improvement to make sure everything aligns with the business requirements.

Capability measure	Level 0	Awareness Level 1	Enablement Level 2	Embedding Level 3
What tools are you using to measure the health of your technical environment?	No visibility	The data types are collected, diagnosis is not readily available.	Existing log traces exist and diagnosis is not easily done.	Logs going to Observe, and logs going to Logstash with dependencies shown.
What does your observability process look like? Do you know if your system is healthy?	Logs going to Observe.	Logs going to Observe.	Logs going to Logstash with dependencies shown.	Detailed logs available and onboarded to obstack.
How do you define and measure your SLIs and SLOs against your requirements?	No level of measure	Service Level SLIs and SLOs are having some issues from for custom.	SLIs and SLOs are generated and reported.	Well defined SLIs and SLOs based on CUs.



Observability maturity self assessment

NOTE

Each level is a different pass of the iterative process. Level 3 demands constant evaluation of the alerts, SLIs, SLOs and continuous improvement to make sure everything aligns with the business requirements.

[New grid](#)

Name of the people working on this

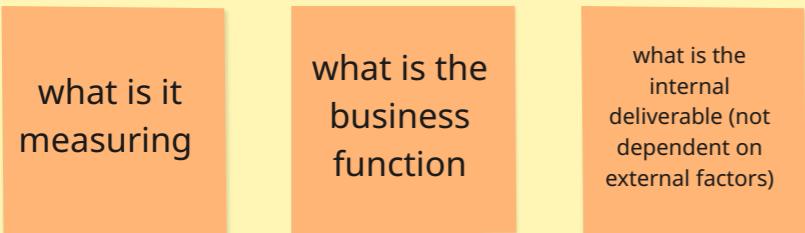
Name of your squad

Capability measure	Current state assessment	Level 0	Current state assessment	Level 1	Current state assessment	Level 2	Current state assessment	Level 3	Current state assessment
What type of observability data exists for your technical service and how easily is it available?	What tools are you using to measure your applications health? Logs or metrics, how can you visualize this? Add details...	The data to support diagnosis of typical failures is not readily collected or available. E.g. No metrics or logs being generated or captured	Add details...	Existing logging, metrics and traces exist to support diagnosing failures, but may require digging or building ad hoc tools. E.g. Logs going to Observe	Add details...	Documentation exists noting the logs, metrics and traces needed to support run books and response to incidents. E.g. Baseline metrics generated and going to Obstack with Dashboard. Logs going to Observe.	Add details...	All observability data for the service is available on real-time dashboards and is sufficient to diagnose a production issue and recover quickly. E.g. Detailed dashboards with dependencies shown. Well defined metrics being generated and onboarded to ObStack. Logs going to Observe.	Add details...
How does your squad find out that your service is broken?	What does your current incident response process look like? How do you know if something is broken? Add details...	Failures of the service do not generate real-time alerts. E.g. No alerts, relying on user complaints.	Add details...	Failures of the service generate alerts but those alerts are delivered to a central team (e.g. Pagerduty), not the engineering team. E.g. Alerts going to FlightDeck	Add details...	Alerts capturing failures are sent to the squad responsible for the failed service who can fix it E.g. Alerts going through PagerDuty based on SLO's.	Add details...	Alerts are received before customer complaints about a degraded/unavailable service. The squad actively maintains these alerts E.g. Different levels of Alerts with critical ones going to PagerDuty Team is using the dashboards to check predicted issues to help mitigate potential issues before they impact customers.	Add details...
How does your squad measure the reliability of your service?	How do you define and measure your applications deliverables against the business expectations? Add details...	The reliability of the service is not measured, or is inferred from for e.g. incident reports or customer complaints E.g. No level of measure	Add details...	A Service Level Indicator (SLI) measuring the service is published with long term data collected E.g. Having some sort of indicator that defines when an alert is triggered.	Add details...	Reliability of the service is measured by exercising at least 1 core business function with an Service Level Objective (SLO) target established E.g. A first pass defined CUJ. Mix of generalised SLI/SLOs and defined SLI/SLOs	Add details...	Service Level Objective (SLO) have coverage of most critical business functions with SLOs target governance in line with business requirements. The Business is regularly providing feedback in SLO governance E.g. Well defined SLI/SLOs based on CUJs. Defined and iterate on CUJs Constant revision of alert thresholds, SLIs,SLOs	Add details...

Braking down SLI and SLO creation

What makes a good SLO

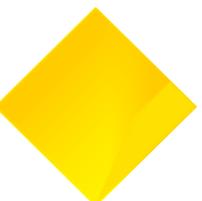
What makes a good SLI



Service Reliability:

- SLIs & SLOs
- Observations

April 2025



Commonwealth
Bank

Service Reliability: SLIs & SLOs

Source: [The Google SRE Book](#)

SLIs: Service Level Indicators

“An SLI is a **service level indicator**, a carefully defined quantitative measure of some aspect of the level of service that is provided.”

SLOs: Service Level Objectives

“An SLO is a **service level objective**, a target value or range of values for a service level that is measured by an SLI.”



Commonwealth
Bank

Service Reliability: SLIs & SLOs

**SLIs: Service Level Indicators =
Reliability Metrics**

An SLI is a ~~service level indicator~~: a carefully defined quantitative measure of some aspect of the level of service that is provided.

**SLOs: Service Level Objectives =
Reliability Targets**

An SLO is a ~~service level objective~~: a target value or range of values for a service level that is measured by an SLI.

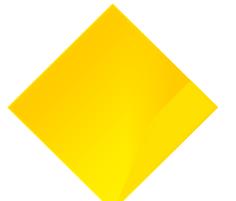


Commonwealth
Bank

SLIs and SLOs at CBA (part 1)

Problem statement; why are we discussing these?

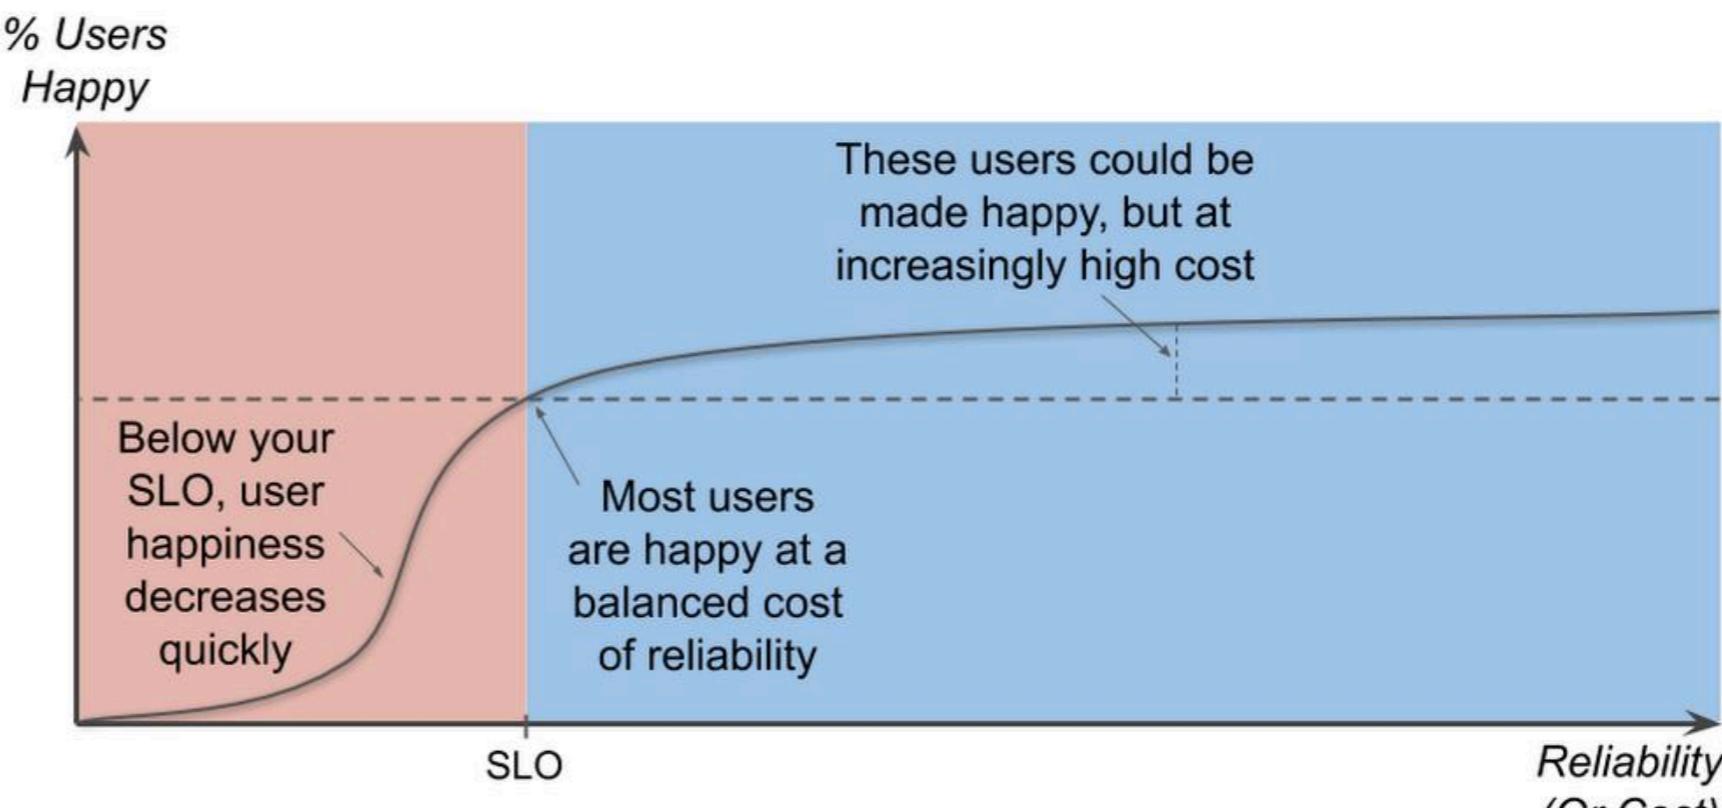
- *CBA requires a uniform way to define, measure and track the reliability of critical services.*
- *Service Level Objectives (SLOs) are a foundation of SRE and provide a standardised way to measure service reliability.*
- *CBA is broadly adopting SLOs as it transitions to a DevSecOps model.*
- In Q1/25, all BU/SUs have an OKR to create fully end-to-end SLOs for at least one of their significant services. The expectation is to build on this experience to develop SLOs for all significant services.
- Central SRE has [published four levels \(0 - 3\) of SLO maturity](#):
 - > CBA has a requirement for all ‘TCF Foundational and Very High Critical services’ to reach and sustain SLO maturity level 2.
- CBA’s Central SRE team is providing guidance on adopting SLOs, with tooling for defining, tracking and visualising them.



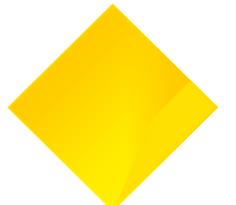
Service reliability

We don't define reliability; critical user journeys do!

- It's important to recognise that it is our users that define what's an acceptable level of service 'reliability'.
- A level of reliability that's too low won't meet user expectations and 'happiness' diminishes rapidly.
- A level of reliability that's too high is expensive, and the marginal gains in user 'happiness' don't justify the extra cost.
- *Well defined SLOs are a balance between meeting user expectations and the costs to implement and sustain a level of reliability.*



Source: [Google Cloud Blog](#)



Reliability levels

Availability is usually expressed as a percentage of uptime in a given period

Reliability Level	Permitted Total Downtime		
	<i>Per year</i>	<i>Per quarter</i>	<i>Per 30 days</i>
90%	36.5 days	9 days	3 days
95%	18.25 days	4.5 days	1.5 days
99%	3.65 days	21.6 hours	7.2 hours
99.5%	1.83 days	10.8 hours	3.6 hours
99.9%	8.76 hours	2.16 hours	43.2 minutes
99.95%	4.38 hours	1.08 hours	21.6 minutes
99.99%	52.6 minutes	12.96 minutes	4.32 minutes
99.999%	5.26 minutes	1.30 minutes	25.9 seconds

← 863ms per day!

Source: [The Google SRE Book](#)

But it's not quite that simple

Reliability Level	Allowed Unreliability Window		
	Per year	Per quarter	Per 30 days
90%	36.5 days	9 days	3 days
95%	18.25 days	4.5 days	1.5 days
99%	3.65 days	21.6 hours	7.2 hours
99.5%	1.83 days	10.8 hours	3.6 hours
99.9%	8.76 hours	2.16 hours	43.2 minutes
99.95%	4.38 hours	1.08 hours	21.6 minutes
99.99%	52.6 minutes	12.96 minutes	4.32 minutes
99.999%	5.26 minutes	1.30 minutes	25.9 seconds

Error Rate	Allowed Duration
100%	21.6 minutes
10%	3.6 hours
1%	36 hours
0.1%	15 days
< 0.05%	All month

Source: [The Google SRE Book](#)

“100% is the wrong reliability target for basically everything.”

Benjamin Treynor Sloss, Vice President of 24x7 Engineering, Google

- 100% reliability is prohibitively expensive – Google’s ‘[Enterprise Roadmap to SRE](#)’ suggests that *every extra ‘9’ in reliability will incur a 10x increase in cost*.

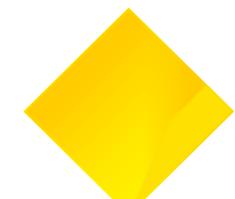
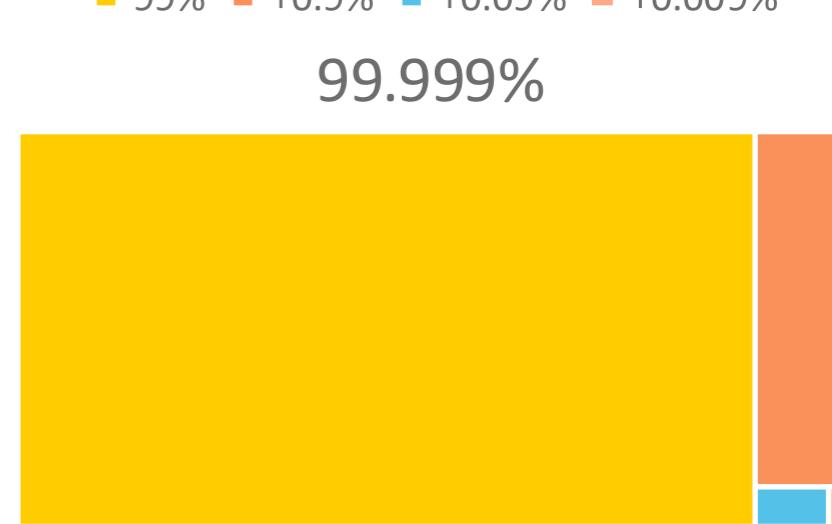
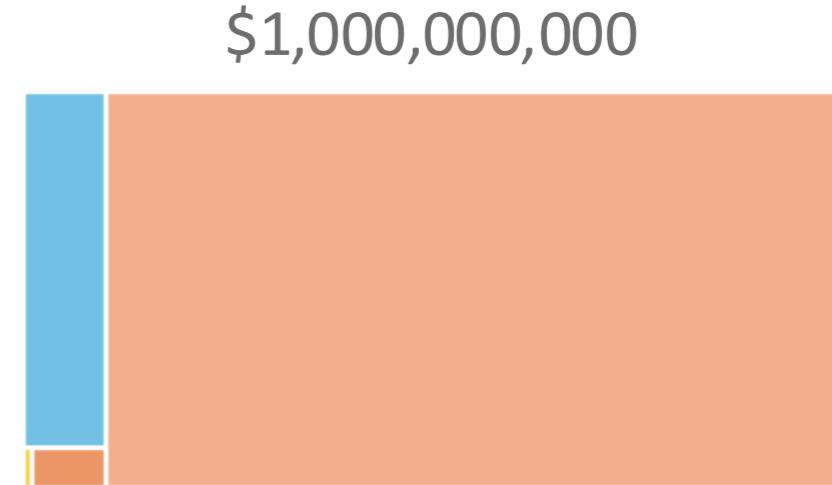
99% @ \$1,000,000 → 99.9% @ \$10,000,000 → 99.99% @ \$100,000,000 → 99.999% @ \$1,000,000,000

- If your SLO is aligned with customer expectations, 100% is simply not a reasonable goal.

- Even if you could achieve 100% reliability within *your system*, your customers would likely not experience *overall* 100% reliability.

- If you do manage to create an experience that is 100% reliable, and want to maintain that level of reliability, you can never update or improve your service.

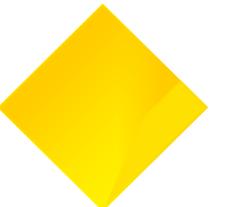
- An SLO of 100% means you only have time to be reactive.



SLO Process Overview

How to SLO? Start simple, start small. Rinse and repeat.

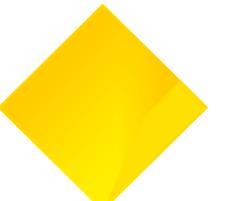
1. Identify critical user journeys (CUJs) and order them by business impact.
2. For each CUJ, determine which metrics to use as service-level indicators (SLIs) to most accurately track the user experience.
3. Determine service-level objective (SLO) target goals and the SLO measurement period.
4. Create SLI and SLO consoles.
5. Create SLO based alerts.



Critical User Journeys (CUJs)

The first step in SLO creation is listing *critical user journeys* for your service.

- A *critical user journey (CUJ)* describes a set of interactions a user/consumer has with a service to achieve some result.
- Let's imagine a web-based shopping site. A few of the actions our customers will be taking when they use the service:
 - > Browse products for sale
 - > Add items into their shopping cart
 - > Check-out and purchase their selected items
- Each of these is a *critical user journey* for this service, and the reliability of each should be measured and tracked.



Service Level Indicators (SLIs)

Metrics that measure some aspect of CUJ service reliability

We recommend treating an SLI as the ratio of two numbers: *The number of good events divided by the total number of events.*

For example:

SLI Type	SLI Definition
Availability	Number of successful HTTP requests / total HTTP requests
Latency	Number of successful RPC requests that complete in < 100ms / total RPC requests
Quality	Number of search results that used the entire corpus / total number of search results, including those that degraded
Freshness	Number of “stock check count” requests that used stock data fresher than 10 minutes / total number of stock check requests

SLIs range from 0% to 100%, where 0% means nothing works, and 100% means nothing is broken.

System and infrastructure metrics are useful, but these are *not* SLIs and rarely have any direct bearing on user experience.

Service Level Objectives (SLOs)

The SLOs for your service are targets for one or more SLIs, aggregated over a specified time

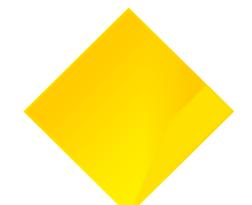
SLI definition		SLO over a four-week rolling window
Availability	<p>The proportion of successful HTTP requests, as measured from the load balancer.</p> <p>Any HTTP status other than 500–599 is considered successful.</p>	99%
Latency	<p>The proportion of sufficiently fast requests, as measured from the load balancer.</p> <p>“Sufficiently fast” is defined as < 200ms, or < 1,000ms.</p>	90% of requests < 200ms 99% of requests < 1,000ms
Correctness	<p>The proportion of records injected into the state table by a prober that result in the correct data being read from the output table.</p> <p>A prober injects synthetic data, with known correct outcomes, and exports a success metric.</p>	99.99% of records injected by the prober result in the correct output

Please do not be tempted to ever use the arithmetic mean in an SLI or SLO. The ‘average’ can be heavily skewed and will often misrepresent your user’s true experience. The 50th percentile, or median, is often a much better measurement.

Future Opportunities

Beyond SLIs and SLOs...

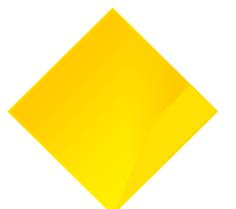
- Single pane of glass view of all critical CBA services and their reliability.
- Consumers of SLOs will include a wider audience than just service owners; SLOs help to reduce silos and provide a common language with which to discuss service reliability across the entire organization.
- SLIs and SLOs will inform data driven decision making: Monitoring, alerting and emergency response, demand forecasting and capacity planning, change management decisions, engineering prioritization.
- A good MTTR is directly related to how quickly you can detect and identify a problem's root cause (the mean time to detect, or MTTD). The longer it takes to identify a problem, the longer it will take you to restore the system to full operation. SLOs will reduce MTTR by shortening MTTD and focusing investigations on a particular CUJ: Targeted, automated health signals.
- Comprehensive SLIs and SLOs are an important foundation piece in the transition to broader DevSecOps & SRE best practices.



SLOs @ CBA: Observations

Some Observability Misunderstandings

1	Obstack is not ready for Production use.
	Obstack has been GA for over twelve months! It's <i>the place to define your SLOs.</i>
2	You cannot put <i>any</i> metrics into Obstack that are not related to CUJs and SLOs.
	We <i>really</i> want you to monitor your services with SLOs, but we won't question up to 20,000 timeseries in total, per application CI. But be wary of cardinality.
3	You cannot monitor your service using logs.
	We <i>really</i> want you to monitor your services with SLOs, and logs should ideally not be your sole source of observability, but they do have merit for certain events: <ul style="list-style-type: none">• High frequency, low cardinality events are best stored as <i>metrics</i>.• However, <i>low frequency, high cardinality events are best kept as logs</i>.
4	You cannot onboard to PagerDuty without CUJs and SLOs.
	We <i>really</i> want you to monitor your services with SLOs, but there is minimal scrutiny of PagerDuty onboarding PRs, and <i>none</i> based on the existence of SLOs.



Cardinality

Thanks to Google Gemini...

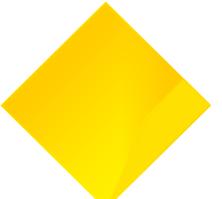
In the context of monitoring and observability, metric cardinality refers to the number of unique timeseries generated by a metric and its associated labels (dimensions). Essentially, it's the count of distinct combinations of metric name and labels. A metric with high cardinality has many unique combinations, while a metric with low cardinality has fewer.

Timeseries:

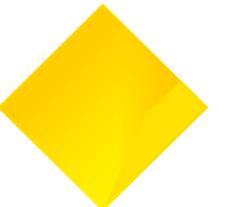
- A time series is a sequence of data points collected over time, with each point representing a specific combination of metric name, labels and values.

Example:

- Imagine a metric called **http_requests_total**. With just the label **status_code**, a metric with a limited number of status codes (e.g., **200, 404, 500**) would have low cardinality. However, if you add labels like **user_id, region**, and **url**, the number of unique combinations would increase dramatically, resulting in high cardinality.



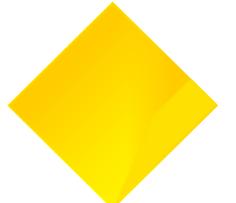
Note to Self: Deep Breath



Monitor for *Symptoms*, not *Causes*

CPU is not an SLI

```
- alert: something-app-prod-ECS-HighCPU-alert
  expr: (aws_ecs_containerinsights_cpu_utilized_average{account_id="3602611xxxxx", dimension_ClusterName="p-s-something-app-ecscluster"} /
         aws_ecs_containerinsights_cpu_reserved_average{account_id="3602611xxxxx", dimension_ClusterName="p-s-something-app-ecscluster"}) * 100 > 70
  for: 10m
  labels:
    cmdb_id: CI0001xxxxx
    severity: warning
  annotations:
    description: Alert - High CPU utilization on Something App ECS.
    summary: High CPU Utilization Detected on Something App ECS.
    message: CPU utilization on ECS has exceeded 70%. Immediate investigation is needed to prevent service degradation or outage .
  runbook_url: https://commbank.atlassian.net/wiki/spaces/Something/pages/12345/Runbook+for+Something+App+Tier+Alerts
```



Monitor for *Symptoms*, not *Causes*

CPU still isn't an SLI

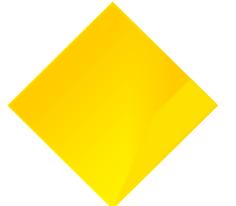
```
- alert: CpuUtilizationMaximum_PrdEgress
  expr: |-
    aws_ec2_cputilization_maximum{account_id="586150xxxxxx"} > 60
  for: 15m
  labels:
    alerter: blackbox
    severity: critical
    cmdb_id: CI0137xxxxx
  annotations:
    message: PRD Egress CPU Utilization is High (Greater than 60 percent) in last 15 minutes.
    runbook_url: "https://commbank.atlassian.net/wiki/spaces/something/pages/921315964/Runbook+-+Workflow+Hosting"
  description: |-  
    PRD Egress High CPU Utilization
```



Monitor for *Symptoms*, not *Causes*

Filesystem

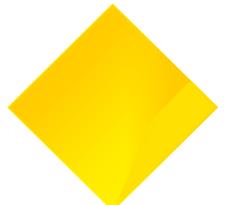
```
- alert: Something-EC2-HighDisk-alert
  expr: (aws_cwagent_disk_used_percent_average{account_id=~"xxxxxxxxxx", region=~"ap-southeast-2"}) > 80
  for: 30m
  labels:
    cmdb_id: CI0001xxxxx
    severity: warning
  annotations:
    description: Alert - High Disk utilization on AWS EC2.
    summary: High Disk Utilization Detected on AWS EC2.
    message: Disk utilization on EC2 has exceeded 80%. Immediate investigation is needed to prevent service degradation or outage .
    runbook_url: https://commbank.atlassian.net/wiki/spaces/Something/pages/8117xxxxx/Runbook+Something+High+Memory+on+all+EC2+instances
```



Monitor for *Symptoms*, not *Causes*

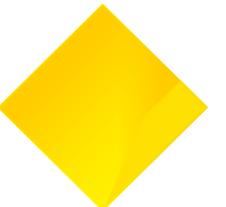
- Monitoring and alerting on system/platform metrics rarely gives any indication as to service health.
- Naïve threshold-based metric alerting is fragile, noisy and misdirected.
- Instead, we need to continuously monitor service health and functionality:
 - > Alert on the *symptoms* of unreliability.
 - > Do not look for *causes* and assume unreliability.
- Platform and system level metrics have their uses in problem diagnosis, capacity planning and trend analysis.
- They are not, however, a good proxy for measuring *system reliability*.

Q: *WHY ARE WE STILL DOING THIS? I have opinions...*



If 100% is wrong, can we please have 99.999%?

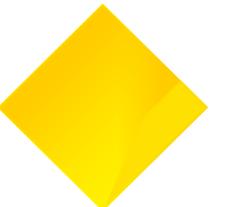
Users may expect at least 99.999% of DNS queries shall be successfully resolved by vhdns. Method: Probe existing records periodically using Blackbox prober. Consider it as successful as long as response != SERVFAIL.



If 100% is wrong, can we please have 99.999%?

Users may expect at least 99.999% of DNS queries shall be successfully resolved by vhdns. Method: Probe existing records periodically using Blackbox prober. Consider it as successful as long as response != SERVFAIL.

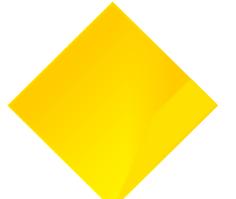
- Be careful what you wish for!
 - > Recall that 99.999% only offers 0.86s of downtime per day; 5.26 total minutes per year.
 - > One failed query from 100,000 is incredibly brittle.



If 100% is wrong, can we please have 99.999%?

Users may expect at least 99.999% of DNS queries shall be successfully resolved by vhdns. Method: Probe existing records periodically using Blackbox prober. Consider it as successful as long as response != SERVFAIL.

- Be careful what you wish for!
 - > Recall that 99.999% only offers 0.86s of downtime per day; 5.26 total minutes per year.
 - > One failed query from 100,000 is incredibly brittle.
- It's not possible to measure 99.999% via a probe. At a minimum, the probe would need to execute every 0.43s to have sufficient samples to measure this (see next slide).
- See <https://uptime.is/99.999>
- Obstack and the Blackbox Exporter don't provide 99.999% availability, and neither does the platform and network that they run on!
- Assuming significant levels of DNS queries, exposing query counters at a loadbalancer or from the binary *may* give sufficient data.





44,100 Hz

https://en.wikipedia.org/wiki/Nyquist-Shannon_sampling_theorem

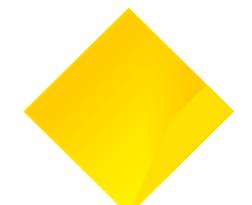
The Nyquist Shannon sampling theorem states that to accurately record an event of duration 'n' you need to sample at a rate of at least ' $n \div 2$ '.

- > A 99.5% SLO over a day allows 432s of downtime, requiring a minimum sample rate of 216s. ✓
- > A 99.9% SLO over a day allows 86.4s of downtime, requiring a minimum sample rate of 43.2s. ✓
- > A 99.95% SLO over a day allows 43.2s of downtime, requiring a minimum sample rate of 21.6s. ?
- > A 99.99% SLO over a day allows 8.64s of downtime, requiring a minimum sample rate of 4.32s. ✗
- > A 99.999% SLO over a day allows 0.864s of downtime, requiring a minimum sample rate of 0.432s. ✗

When setting your SLOs, be aware of the data granularity required to support it!

Unless you can probe exceedingly fast or have a service with high transactions per second and a means of exporting accurate statistics, you will very likely not have sufficient data to calculate your SLO adherence.

Or more simply, *you will miss events.*

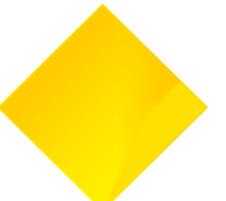


My SLO only applies when my system is being used!

- SLOs, and the Sloth framework utilised within Obstack, are designed to track continuous 24-hour availability.
- Many CBA systems have historically only offered ‘business hours’ availability.
- Many more have well defined periods of scheduled, sustained unavailability (maintenance).

Q: Should SLOs only apply during documented periods of expected, scheduled availability?

- The SRE purist answer is that SLOs should be continuous and inclusive of maintenance events, etc.
 - The pragmatic answer might be that today CBA does have systems that operate within timebound limits, and SLOs at CBA might need to accommodate these (for now)?
- > SPOILER: Obstack and Sloth have *no* support for non-24-hour SLOs.



Thank-you!

Will Charles, Principal Engineer, Core SRE

will.charles@cba.com.au

<http://edu.sre.cba/contact/contact-us> and <http://chat.sre.cba/>

